

Insurance claim prediction

March 19, 2024

1 Insurance claim prediction

1.1 Problem Statement :

Dataset is about Vehicle insurance claim prediction in which we have two datasets one is train and another is test data. The challenge is to predict Workers Compensation claims using the realistic data .

1.2 Independent/Predictor Variables

- ClaimNumber: Unique policy identifier
- DateTimeOfAccident: Date and time of accident
- DateReported: Date that accident was reported
- Age: Age of worker
- Gender: Gender of worker
- MaritalStatus: Martial status of worker. (M)arried, (S)ingle, (U)unknown.
- DependentChildren: The number of dependent children
- DependentsOther: The number of dependants excluding children
- WeeklyWages: Total weekly wage
- PartTimeFullTime: Binary (P) or (F)
- HoursWorkedPerWeek: Total hours worked per week
- DaysWorkedPerWeek: Number of days worked per week
- ClaimDescription: Free text description of the claim
- InitialIncurredClaimCost: Initial estimate by the insurer of the claim cost

1.3 Dependent/Target Variable

- UltimateIncurredClaimCost: Total claims payments by the insurance company.
- The Target variable 'Ultimate Incurred Claim Cost' is a continuous variable which accounts for the total claim payment by the insurance company

2 Importing Libraries

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
import warnings
warnings.filterwarnings("ignore")
```

3 Data Preparation

```
[2]: #collecting data
df = pd.read_csv('train_SJC.csv')
df.head()
```

```
[2]:      Unnamed: 0      Unnamed: 1      DateReported Unnamed: 3 \
0  ClaimNumber  DateTimeOfAccident      NaN      Age
1    WC8205482  2002-04-09T07:00:00Z  2002-07-05T00:00:00Z      48
2    WC6922469  1999-01-07T11:00:00Z  1999-01-20T00:00:00Z      43
3    WC5442654  1996-03-25T00:00:00Z  1996-04-14T00:00:00Z      30
4    WC9796897  2005-06-22T13:00:00Z  2005-07-22T00:00:00Z      41

      Unnamed: 4      Unnamed: 5      Unnamed: 6  DependentsOther  Unnamed: 8 \
0      Gender  MaritalStatus  DependentChildren      NaN  WeeklyWages
1          M          M          0          0.0          500
2          F          M          0          0.0          509.34
3          M          U          0          0.0          709.1
4          M          S          0          0.0          555.46

      Unnamed: 9      Unnamed: 10  DaysWorkedPerWeek \
0  PartTimeFullTime  HoursWorkedPerWeek      NaN
1          F          38          5.0
2          F          37.5          5.0
3          F          38          5.0
4          F          38          5.0

      Unnamed: 12 \
0          ClaimDescription
1  LIFTING TYRE INJURY TO RIGHT ARM AND WRIST INJURY
2  STEPPED AROUND CRATES AND TRUCK TRAY FRACTURE ...
3          CUT ON SHARP EDGE CUT LEFT THUMB
4          DIGGING LOWER BACK LOWER BACK STRAIN

      Unnamed: 13      Unnamed: 14
0  InitialIncurredCalimsCost  UltimateIncurredClaimCost
1          1500          4748.203388
2          5500          6326.285819
3          1700          2293.949087
4         15000          17786.48717
```

```
[3]: #Renaming the column names in train data as it is "unnamed"
```

```
df=df.rename(columns={"Unnamed: 0":"ClaimNumber","Unnamed: 1":
↳"DateTimeOfAccident","Unnamed: 3":"Age","Unnamed: 4":"Gender",
"Unnamed: 5":"MaritalStatus","Unnamed: 6":
↳"DependentChildren","Unnamed: 8":"WeeklyWages",
"Unnamed: 9":"PartTimeFullTime","Unnamed: 10":
↳"HoursWorkedPerWeek","Unnamed: 12":"ClaimDescription",
"Unnamed: 13":"InitialIncurredCalimsCost","Unnamed: 14":
↳'UltimateIncurredClaimCost'},inplace=False)
```

```
[4]: df.head()
```

```
[4]:
```

	ClaimNumber	DateTimeOfAccident	DateReported	Age	Gender	\
0	ClaimNumber	DateTimeOfAccident		NaN	Age	Gender
1	WC8205482	2002-04-09T07:00:00Z	2002-07-05T00:00:00Z	48		M
2	WC6922469	1999-01-07T11:00:00Z	1999-01-20T00:00:00Z	43		F
3	WC5442654	1996-03-25T00:00:00Z	1996-04-14T00:00:00Z	30		M
4	WC9796897	2005-06-22T13:00:00Z	2005-07-22T00:00:00Z	41		M

	MaritalStatus	DependentChildren	DependentsOther	WeeklyWages	\
0	MaritalStatus	DependentChildren		NaN	WeeklyWages
1	M	0	0.0		500
2	M	0	0.0		509.34
3	U	0	0.0		709.1
4	S	0	0.0		555.46

	PartTimeFullTime	HoursWorkedPerWeek	DaysWorkedPerWeek	\
0	PartTimeFullTime	HoursWorkedPerWeek		NaN
1	F	38		5.0
2	F	37.5		5.0
3	F	38		5.0
4	F	38		5.0

	ClaimDescription	\
0	ClaimDescription	
1	LIFTING TYRE INJURY TO RIGHT ARM AND WRIST INJURY	
2	STEPPED AROUND CRATES AND TRUCK TRAY FRACTURE ...	
3	CUT ON SHARP EDGE CUT LEFT THUMB	
4	DIGGING LOWER BACK LOWER BACK STRAIN	

	InitialIncurredCalimsCost	UltimateIncurredClaimCost
0	InitialIncurredCalimsCost	UltimateIncurredClaimCost
1	1500	4748.203388
2	5500	6326.285819
3	1700	2293.949087
4	15000	17786.48717

```
[5]: df=df.drop(df.index[0])
df.head()
```

```
[5]: ClaimNumber      DateTimeOfAccident      DateReported Age Gender \
1      WC8205482      2002-04-09T07:00:00Z      2002-07-05T00:00:00Z  48      M
2      WC6922469      1999-01-07T11:00:00Z      1999-01-20T00:00:00Z  43      F
3      WC5442654      1996-03-25T00:00:00Z      1996-04-14T00:00:00Z  30      M
4      WC9796897      2005-06-22T13:00:00Z      2005-07-22T00:00:00Z  41      M
5      WC2603726      1990-08-29T08:00:00Z      1990-09-27T00:00:00Z  36      M
```

```
MaritalStatus DependentChildren DependentsOther WeeklyWages \
1      M      0      0.0      500
2      M      0      0.0      509.34
3      U      0      0.0      709.1
4      S      0      0.0      555.46
5      M      0      0.0      377.1
```

```
PartTimeFullTime HoursWorkedPerWeek DaysWorkedPerWeek \
1      F      38      5.0
2      F      37.5      5.0
3      F      38      5.0
4      F      38      5.0
5      F      38      5.0
```

```
ClaimDescription \
1      LIFTING TYRE INJURY TO RIGHT ARM AND WRIST INJURY
2      STEPPED AROUND CRATES AND TRUCK TRAY FRACTURE ...
3      CUT ON SHARP EDGE CUT LEFT THUMB
4      DIGGING LOWER BACK LOWER BACK STRAIN
5      REACHING ABOVE SHOULDER LEVEL ACUTE MUSCLE STR...
```

```
InitialIncurredCalimsCost UltimateIncurredClaimCost
1      1500      4748.203388
2      5500      6326.285819
3      1700      2293.949087
4      15000      17786.48717
5      2800      4014.002925
```

```
[6]: #for training data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 36176 entries, 1 to 36176
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   ClaimNumber         36176 non-null  object
1   DateTimeOfAccident  36176 non-null  object
```

```

2   DateReported          36176 non-null object
3   Age                   36176 non-null object
4   Gender                36176 non-null object
5   MaritalStatus         36154 non-null object
6   DependentChildren     36176 non-null object
7   DependentsOther       36176 non-null float64
8   WeeklyWages           36120 non-null object
9   PartTimeFullTime      36176 non-null object
10  HoursWorkedPerWeek     36127 non-null object
11  DaysWorkedPerWeek      36176 non-null float64
12  ClaimDescription       36176 non-null object
13  InitialIncurredCalimsCost 36176 non-null object
14  UltimateIncurredClaimCost 36176 non-null object
dtypes: float64(2), object(13)
memory usage: 4.1+ MB

```

```

[7]: #Checking for Categorical Data in train data
df.select_dtypes(exclude=['int64', 'float64']).columns

```

```

[7]: Index(['ClaimNumber', 'DateTimeOfAccident', 'DateReported', 'Age', 'Gender',
        'MaritalStatus', 'DependentChildren', 'WeeklyWages', 'PartTimeFullTime',
        'HoursWorkedPerWeek', 'ClaimDescription', 'InitialIncurredCalimsCost',
        'UltimateIncurredClaimCost'],
        dtype='object')

```

```

[8]: #Changing the data type for some columns in train data

df['Age'] = pd.to_numeric(df['Age'])
df['DependentChildren'] = pd.to_numeric(df['DependentChildren'])
df['DependentsOther'] = pd.to_numeric(df['DependentsOther'])
df['WeeklyWages'] = pd.to_numeric(df['WeeklyWages'])
df['HoursWorkedPerWeek'] = pd.to_numeric(df['HoursWorkedPerWeek'])
df['DaysWorkedPerWeek'] = pd.to_numeric(df['DaysWorkedPerWeek'])
df['InitialIncurredCalimsCost'] = pd.to_numeric(df['InitialIncurredCalimsCost'])
df['UltimateIncurredClaimCost'] = pd.to_numeric(df['UltimateIncurredClaimCost'])
df['DateTimeOfAccident'] = pd.to_datetime(df['DateTimeOfAccident'])
df['DateReported'] = pd.to_datetime(df['DateReported'])

```

```

[9]: #To check if the data type changed or not
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 36176 entries, 1 to 36176
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ClaimNumber                          36176 non-null  object
1   DateTimeOfAccident                   36176 non-null  datetime64[ns, UTC]

```

```

2   DateReported      36176 non-null  datetime64[ns, UTC]
3   Age               36176 non-null  int64
4   Gender            36176 non-null  object
5   MaritalStatus     36154 non-null  object
6   DependentChildren 36176 non-null  int64
7   DependentsOther   36176 non-null  float64
8   WeeklyWages       36120 non-null  float64
9   PartTimeFullTime  36176 non-null  object
10  HoursWorkedPerWeek 36127 non-null  float64
11  DaysWorkedPerWeek  36176 non-null  float64
12  ClaimDescription   36176 non-null  object
13  InitialIncurredCalimsCost 36176 non-null  int64
14  UltimateIncurredClaimCost 36176 non-null  float64
dtypes: datetime64[ns, UTC](2), float64(5), int64(3), object(5)
memory usage: 4.1+ MB

```

```

[10]: #checking the description of train data
df.describe()

```

```

[10]:
count      36176.000000    36176.000000    36176.000000    36120.000000 \
mean       33.795196         0.121296         0.009537         416.471426
std        12.114729         0.525395         0.106163         243.875364
min        13.000000         0.000000         0.000000          1.000000
25%        23.000000         0.000000         0.000000         200.000000
50%        32.000000         0.000000         0.000000         393.365000
75%        43.000000         0.000000         0.000000         500.000000
max        79.000000         9.000000         3.000000        7497.000000

count      36127.000000    36176.000000    36176.000000 \
mean       37.766820         4.905794         7743.593874
std        12.494323         0.547077        18223.698531
min         0.000000         1.000000          1.000000
25%        38.000000         5.000000          700.000000
50%        38.000000         5.000000        2000.000000
75%        40.000000         5.000000        9500.000000
max        640.000000         7.000000       83000.000000

count      3.617600e+04
mean       1.095282e+04
std        3.529614e+04
min        1.218868e+02
25%        9.257424e+02
50%        3.373862e+03
75%        8.186852e+03

```

max 4.027136e+06

```
[11]: #Checking for duplicate values
df.duplicated().sum()
```

[11]: 0

```
[12]: #Checking for null values
df.isnull().sum()
```

```
[12]: ClaimNumber          0
      DateTimeOfAccident   0
      DateReported         0
      Age                  0
      Gender               0
      MaritalStatus        22
      DependentChildren    0
      DependentsOther      0
      WeeklyWages          56
      PartTimeFullTime     0
      HoursWorkedPerWeek   49
      DaysWorkedPerWeek    0
      ClaimDescription     0
      InitialIncurredCalimsCost  0
      UltimateIncurredClaimCost  0
      dtype: int64
```

```
[13]: #Handling null values using mean and mode imputation to treating the missing_
      ↪values
df['WeeklyWages']=df['WeeklyWages'].fillna(df['WeeklyWages'].mean())
df['HoursWorkedPerWeek']=df['HoursWorkedPerWeek'].
      ↪fillna(df['HoursWorkedPerWeek'].mean())
df['MaritalStatus']=df['MaritalStatus'].fillna(df['MaritalStatus'].mode()[0])
```

```
[14]: #To verify if there are any more missing values
df.isnull().sum()
```

```
[14]: ClaimNumber          0
      DateTimeOfAccident   0
      DateReported         0
      Age                  0
      Gender               0
      MaritalStatus        0
      DependentChildren    0
      DependentsOther      0
      WeeklyWages          0
      PartTimeFullTime     0
```

```

HoursWorkedPerWeek      0
DaysWorkedPerWeek       0
ClaimDescription         0
InitialIncurredCalimsCost 0
UltimateIncurredClaimCost 0
dtype: int64

```

No Null values are present

Dividing the data into categorical and numerical data

```

[15]: df_num = df[['Age', 'DependentChildren', 'DependentsOther', 'WeeklyWages',
↳ 'HoursWorkedPerWeek', 'DaysWorkedPerWeek',
        'InitialIncurredCalimsCost', 'UltimateIncurredClaimCost']]

df_cat = df[['ClaimNumber', 'DateTimeOfAccident', 'DateReported', 'Gender',
↳ 'MaritalStatus', 'PartTimeFullTime', 'ClaimDescription']]

```

```
[16]: df['ClaimDescription'].nunique()
```

```
[16]: 20596
```

There are 20596 unique claims made.

```
[17]: df.ClaimNumber.count()
```

```
[17]: 36176
```

The total number of claims filed is 36176.

```
[18]: df['ClaimNumber'].nunique()
```

```
[18]: 29456
```

The total number of claims that were filed is 36176 but number of unique claims are 29456.

Data Transformation - Data Binning

```
[19]: df['Age'].value_counts
```

```

[19]: <bound method IndexOpsMixin.value_counts of 1      48
      2      43
      3      30
      4      41
      5      36
      ..
    36172    20
    36173    35
    36174    52
    36175    28
    36176    29

```


Name: Age, Length: 36176, dtype: int64>

```
[20]: df['Age'].max()
```

```
[20]: 79
```

```
[21]: df['Age'].min()
```

```
[21]: 13
```

```
[22]: df['Age_Bin']=pd.cut(df['Age'],bins=[1,25,50,80] ,  
    ↪labels=['Young','Middle-Age','Old'])  
df['Age_Bin']
```

```
[22]: 1      Middle-Age  
2      Middle-Age  
3      Middle-Age  
4      Middle-Age  
5      Middle-Age  
...  
36172      Young  
36173      Middle-Age  
36174      Old  
36175      Middle-Age  
36176      Middle-Age  
Name: Age_Bin, Length: 36176, dtype: category  
Categories (3, object): ['Young' < 'Middle-Age' < 'Old']
```

```
[23]: df['WeeklyWages'].value_counts
```

```
[23]: <bound method IndexOpsMixin.value_counts of 1      500.00  
2      509.34  
3      709.10  
4      555.46  
5      377.10  
...  
36172      344.16  
36173      1668.83  
36174      204.87  
36175      730.87  
36176      200.00  
Name: WeeklyWages, Length: 36176, dtype: float64>
```

```
[24]: df['WeeklyWages'].max()
```

```
[24]: 7497.0
```

```
[25]: df['WeeklyWages'].min()
```

```
[25]: 1.0
```

```
[26]: df['WeeklyWages_Bin']=pd.  
      ↪cut(df['WeeklyWages'],bins=[0,1000,2000,4000,7000,8000] ,  
      ↪labels=['Low','Below Average','Average Wage','Above Average','High'])  
      df['WeeklyWages_Bin']
```

```
[26]: 1          Low  
      2          Low  
      3          Low  
      4          Low  
      5          Low  
      ...  
      36172       Low  
      36173  Below Average  
      36174       Low  
      36175       Low  
      36176       Low  
      Name: WeeklyWages_Bin, Length: 36176, dtype: category  
      Categories (5, object): ['Low' < 'Below Average' < 'Average Wage' < 'Above  
      Average' < 'High']
```

4 Exploratory Data Analysis :

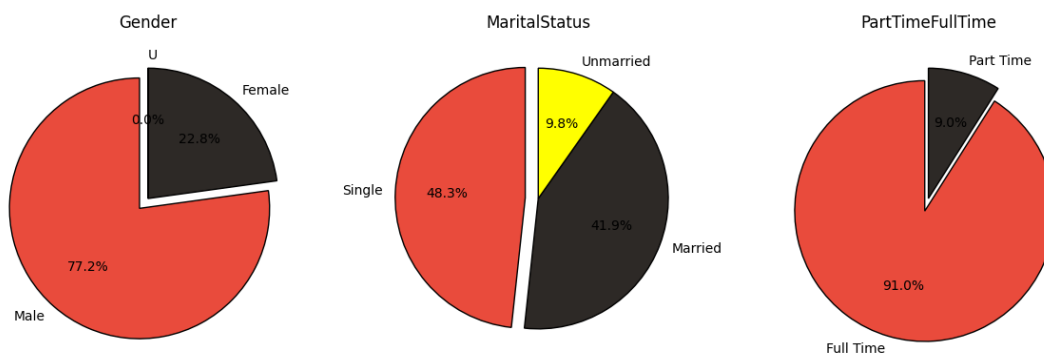
4.1 Univariate Analysis

```
[27]: Gender = df['Gender'].value_counts()  
      Gender = [Gender[0] / sum(Gender) * 100, Gender[1] / sum(Gender) * 100 ,  
      ↪Gender[2] / sum(Gender) * 100]  
  
      MaritalStatus = df['MaritalStatus'].value_counts()  
      MaritalStatus = [MaritalStatus[0] / sum(MaritalStatus) * 100, MaritalStatus[1] /  
      ↪sum(MaritalStatus) * 100 , MaritalStatus[2] / sum(MaritalStatus) * 100]  
  
      PartTimeFullTime = df['PartTimeFullTime'].value_counts()  
      PartTimeFullTime = [PartTimeFullTime[0] / sum(PartTimeFullTime) * 100,  
      ↪PartTimeFullTime[1] / sum(PartTimeFullTime) * 100 ]
```

```
[28]: colors = ['#E94B3C','#2D2926','#FFFF00']  
  
      ax,fig = plt.subplots(nrows = 1,ncols = 3,figsize = (15,15))  
  
      plt.subplot(1,3,1)  
      plt.pie(Gender,labels = ['Male','Female','U'],autopct='%1.1f%%',startangle =  
      ↪90,explode = (0.1,0,0),colors = colors,  
      ↪wedgeprops = {'edgecolor' : 'black','linewidth': 1,'antialiased' : True})  
      plt.title('Gender');
```

```
plt.subplot(1,3,2)
plt.pie(MaritalStatus,labels = ['Single', 'Married','Unmarried'],autopct='%1.
↪1f%%',startangle = 90,explode = (0.1,0,0),colors = colors,
      wedgeprops = {'edgecolor' : 'black','linewidth': 1,'antialiased' : True})
plt.title('MaritalStatus');

plt.subplot(1,3,3)
plt.pie(PartTimeFullTime,labels = ['Full Time','Part Time'],autopct='%1.
↪1f%%',startangle = 90,explode = (0.1,0),colors = colors,
      wedgeprops = {'edgecolor' : 'black','linewidth': 1,'antialiased' : True})
plt.title('PartTimeFullTime');
```



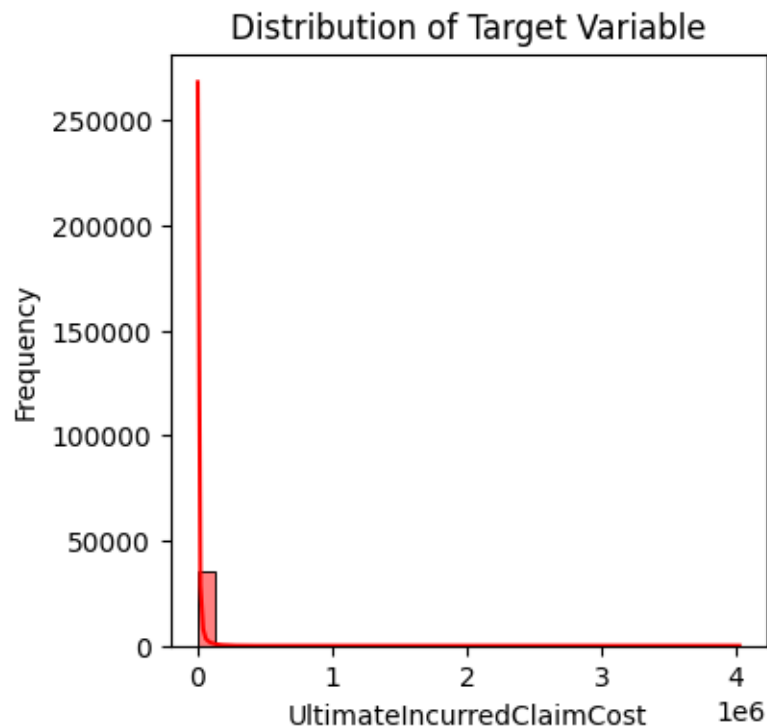
- From the above plot we can see approximately 77% people who claimed for insurance are males while 22% are females.
- We can see approximately 48% people who claimed for insurance are single while 42% are married and a very small percentage of people are unmarried.
- We can see approximately 91% of the people hold full time jobs.

```
[29]: #Checking the target variable
df['UltimateIncurredClaimCost'].describe()
```

```
[29]: count      3.617600e+04
      mean       1.095282e+04
      std        3.529614e+04
      min        1.218868e+02
      25%         9.257424e+02
      50%         3.373862e+03
      75%         8.186852e+03
      max         4.027136e+06
      Name: UltimateIncurredClaimCost, dtype: float64
```

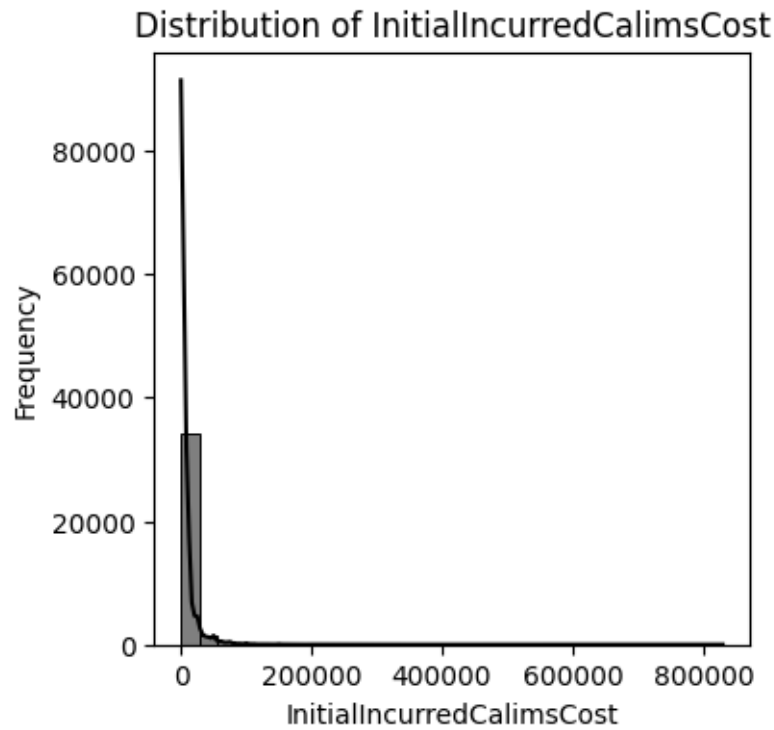
```
[30]: #Checking the skewness of the target variable
plt.figure(figsize=(4,4))
```

```
sns.histplot(data=df, x='UltimateIncurredClaimCost', kde=True,
             ↪bins=30,color='red')
plt.title('Distribution of Target Variable')
plt.xlabel('UltimateIncurredClaimCost')
plt.ylabel('Frequency')
plt.show()
```



- The data for UltimateIncurredClaimCost is right skewed.

```
[31]: plt.figure(figsize=(4,4))
sns.histplot(data=df, x='InitialIncurredCalimsCost', kde=True, bins=30,color=
             ↪'black')
plt.title('Distribution of InitialIncurredCalimsCost')
plt.xlabel('InitialIncurredCalimsCost')
plt.ylabel('Frequency')
plt.show()
```



- The data for InitialIncurredClaimCost is right skewed.

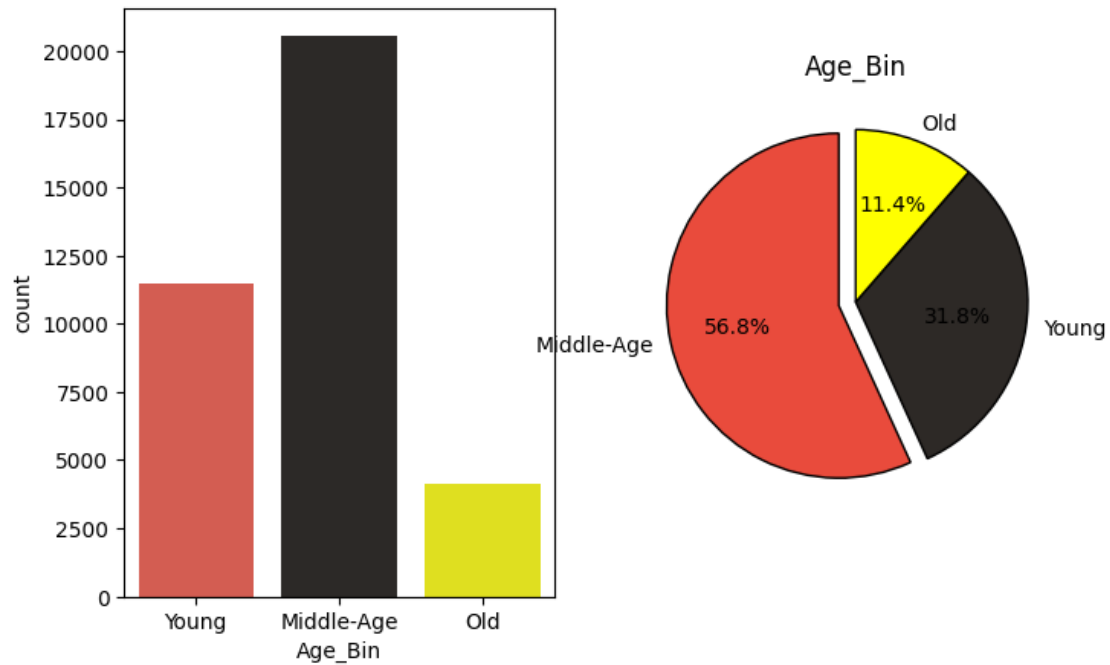
```
[32]: plt.figure(figsize=(8, 6))
sns.histplot(data=df, x='Age', kde=True, bins=30,color='red')
plt.title('Distribution of Age')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```



```
[33]: Age_Bin= df['Age_Bin'].value_counts(normalize=True)*100
ax,fig = plt.subplots(nrows = 1,ncols = 3,figsize = (8,5))

plt.subplot(1,2,1)
sns.countplot(x = 'Age_Bin', data = df,palette=colors)

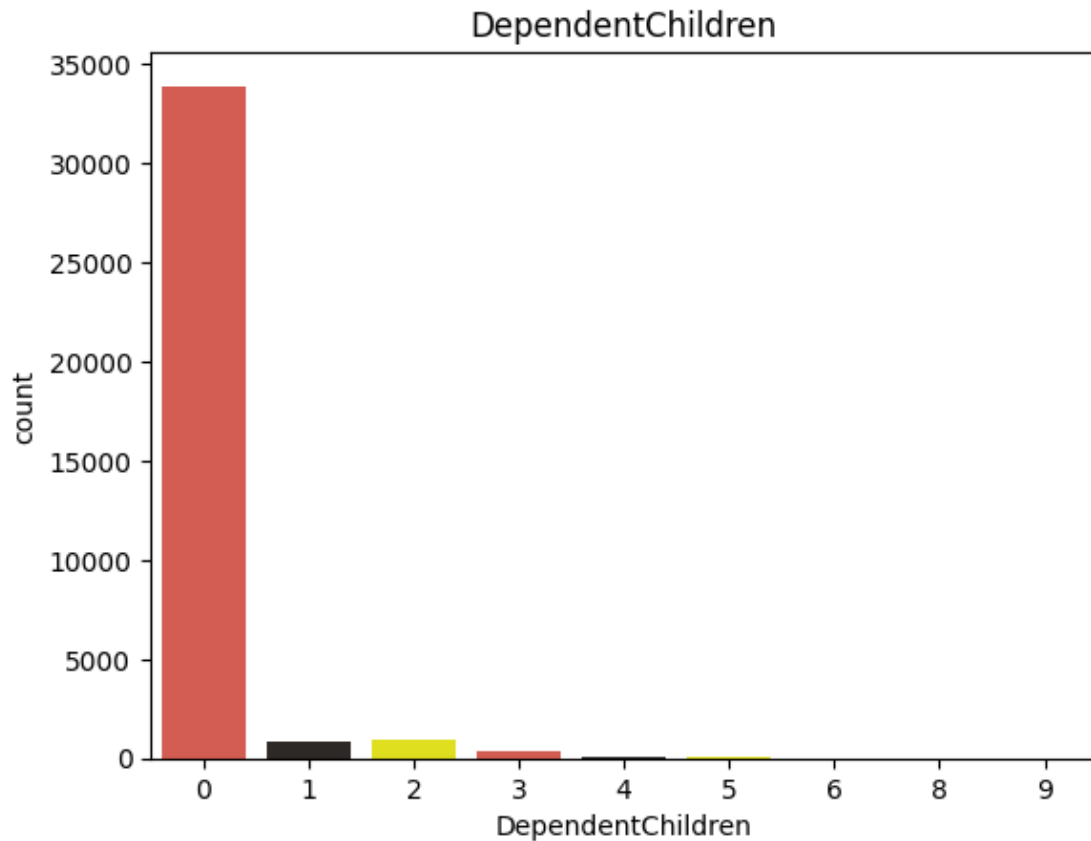
plt.subplot(1,2,2)
plt.pie(Age_Bin,labels = ['Middle-Age','Young','Old'],autopct='%1.
    ↪1f%',startangle = 90,explode = (0.1,0,0),colors = colors,
        wedgeprops = {'edgecolor' : 'black','linewidth': 1,'antialiased' : True})
plt.title('Age_Bin');
```



- From both the plots we can see that the claims are higher from the middle age group (25-40) which is roughly 57%.

```
[34]: sns.countplot(x = 'DependentChildren', data = df,palette=colors)
      plt.title('DependentChildren')
```

```
[34]: Text(0.5, 1.0, 'DependentChildren')
```



```
[35]: df['DependentChildren'].value_counts(normalize=True)*100
```

```
[35]: DependentChildren
```

```
0    93.703008
2     2.554180
1     2.371738
3     0.975785
4     0.284719
5     0.093985
6     0.011057
9     0.002764
8     0.002764
```

```
Name: proportion, dtype: float64
```

- The claims made from people with no children as dependents is high(Almost 94%).

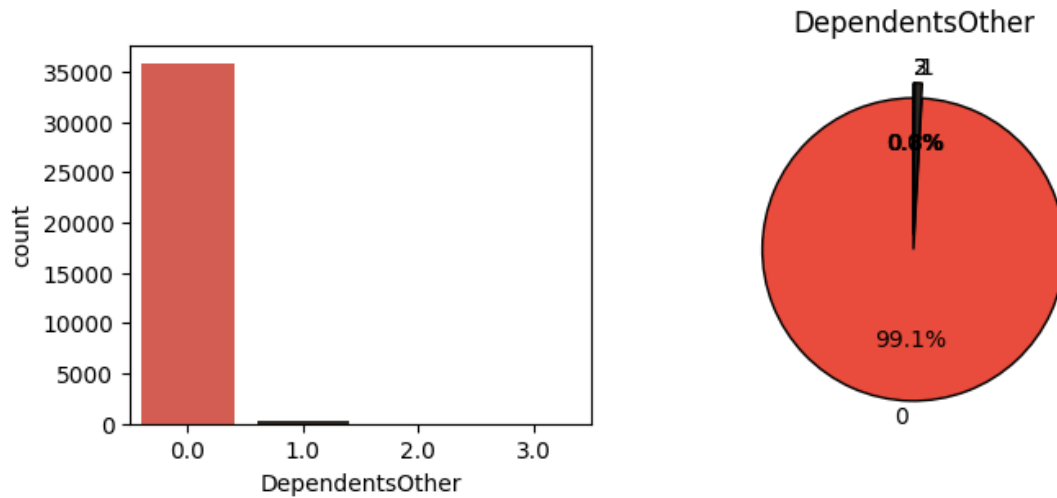
```
[36]: DependentsOther= df['DependentsOther'].value_counts(normalize=True)*100
```

```
ax,fig = plt.subplots(nrows = 1,ncols = 3,figsize = (8,3))
```



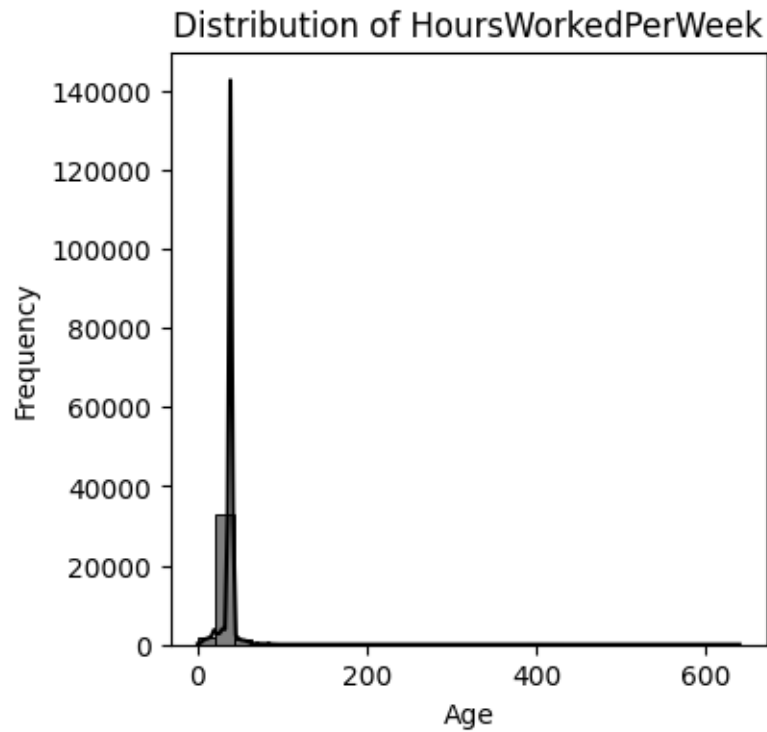
```
plt.subplot(1,2,1)
sns.countplot(x = 'DependentsOther', data = df,palette=colors)

plt.subplot(1,2,2)
plt.pie(DependentsOther,labels = ['0','1','2','3'],autopct='%1.1f%%',startangle=
    ↪ 90,explode = (0.1,0,0,0),colors = colors,
        wedgeprops = {'edgecolor' : 'black','linewidth': 1,'antialiased' : True})
plt.title('DependentsOther');
```



- The claims made from people with no dependents is high (Almost 99%).

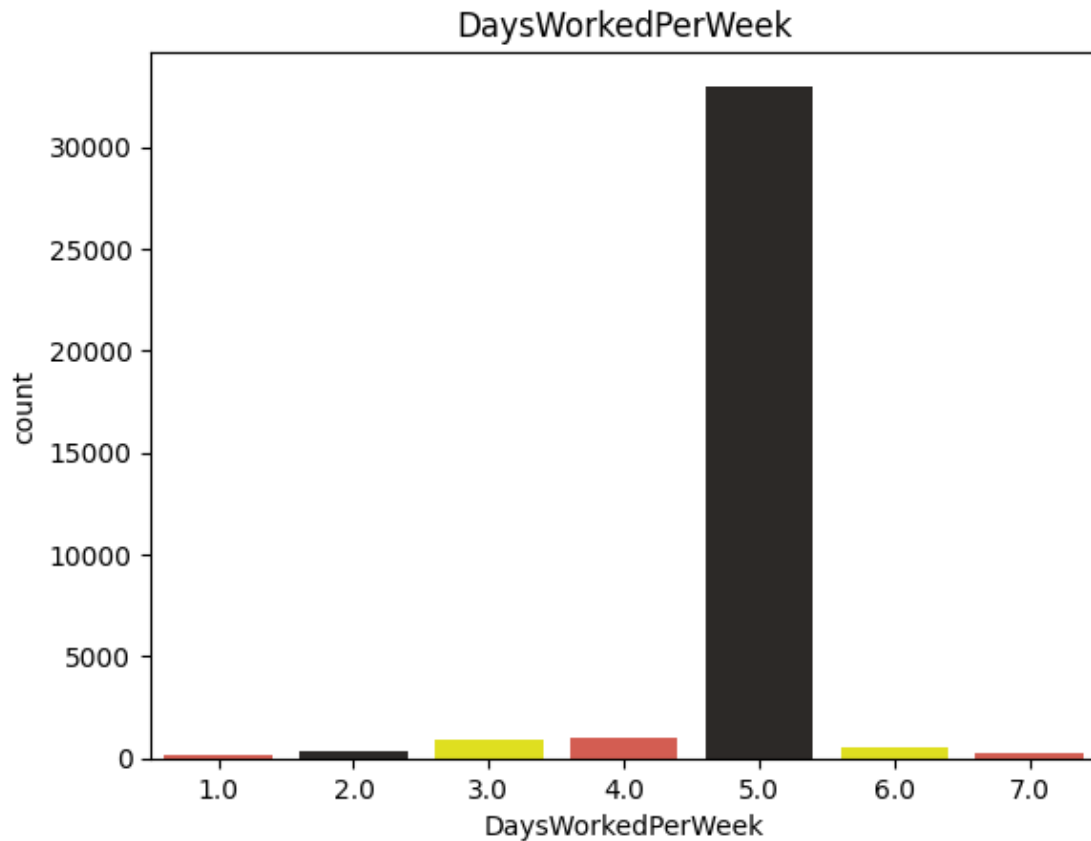
```
[37]: plt.figure(figsize=(4,4))
sns.histplot(data=df, x='HoursWorkedPerWeek', kde=True, bins=30,color='black')
plt.title('Distribution of HoursWorkedPerWeek')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```



- The data for HoursWorkedPerWeek is positively skewed.

```
[38]: sns.countplot(x = 'DaysWorkedPerWeek', data = df , palette=colors)
plt.title('DaysWorkedPerWeek')
```

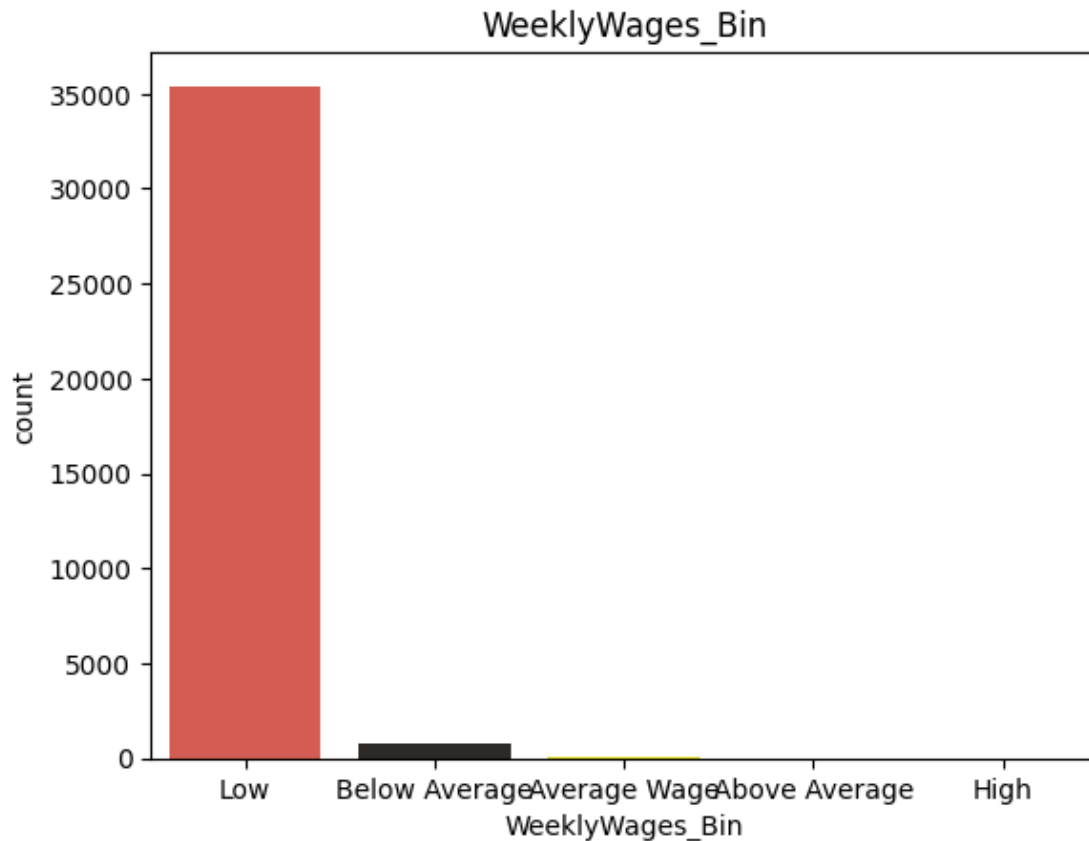
```
[38]: Text(0.5, 1.0, 'DaysWorkedPerWeek')
```



- From the above plot we can see that the people who have claimed for insurance work for 5 days a week.

```
[39]: sns.countplot(x = 'WeeklyWages_Bin', data = df , palette=colors)
plt.title('WeeklyWages_Bin')
```

```
[39]: Text(0.5, 1.0, 'WeeklyWages_Bin')
```



```
[40]: df['WeeklyWages_Bin'].value_counts(normalize=True)*100
```

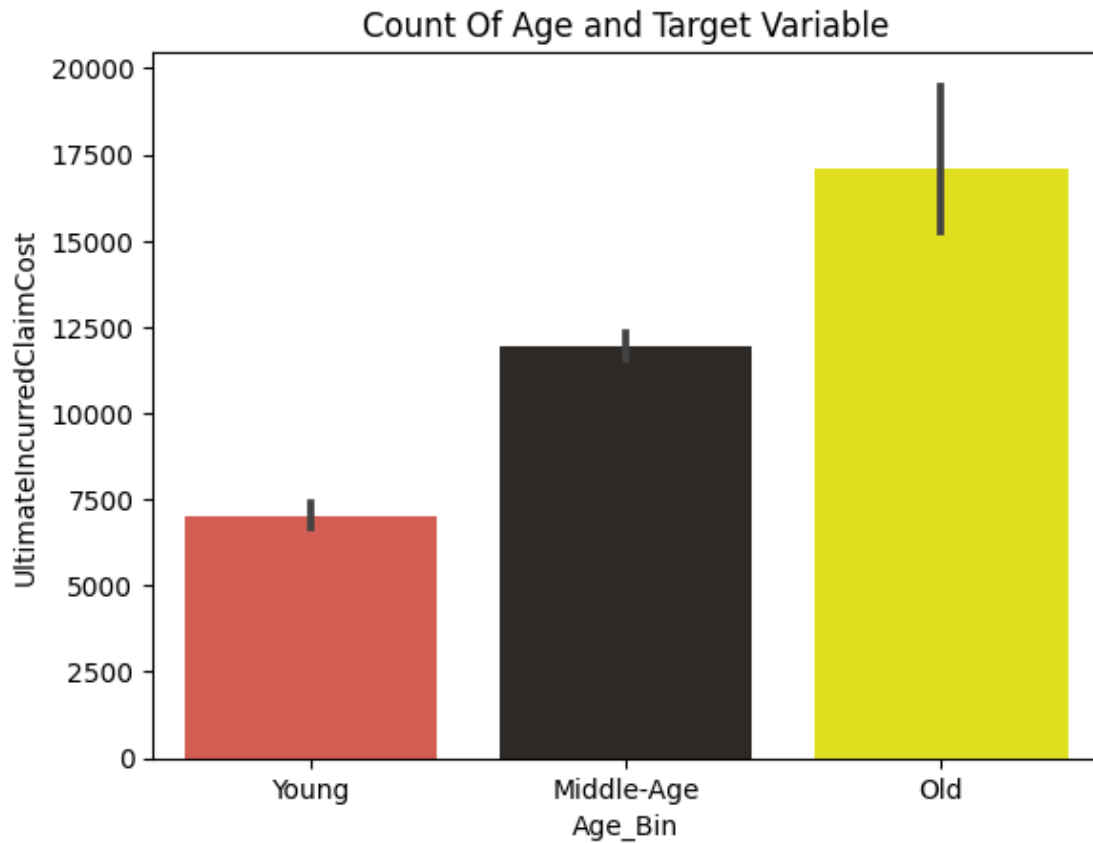
```
[40]: WeeklyWages_Bin
Low          97.910217
Below Average    1.973684
Average Wage    0.105042
Above Average    0.008293
High           0.002764
Name: proportion, dtype: float64
```

- From the above plot we can see that most of the people who claimed for insurance have low wages

4.2 Bivariate analysis

```
[41]: sns.barplot(x='Age_Bin',y='UltimateIncurredClaimCost',data=df , palette=colors)
plt.title('Count Of Age and Target Variable')
```

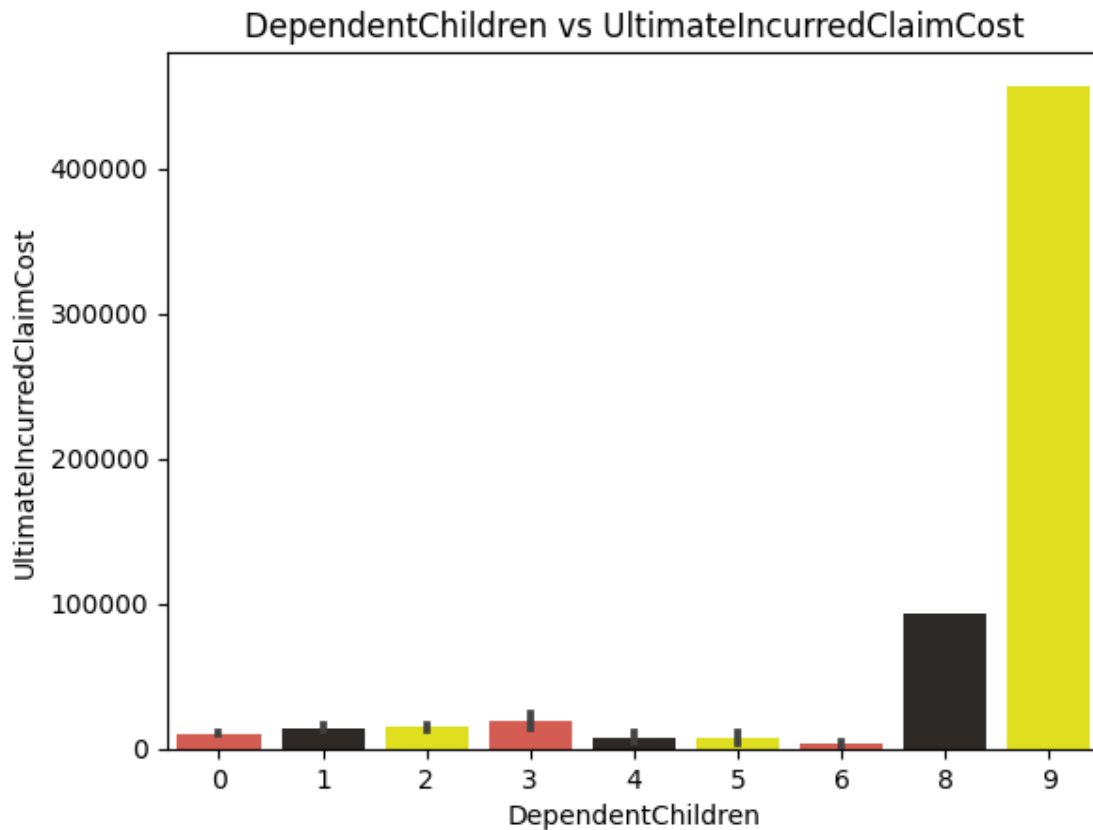
```
[41]: Text(0.5, 1.0, 'Count Of Age and Target Variable')
```



- People who are in the old age group (50-80) got more total claims payments by the insurance company.

```
[42]: sns.barplot(x='DependentChildren',y='UltimateIncurredClaimCost',data=df_
      ↪,palette=colors)
plt.title('DependentChildren vs UltimateIncurredClaimCost')
```

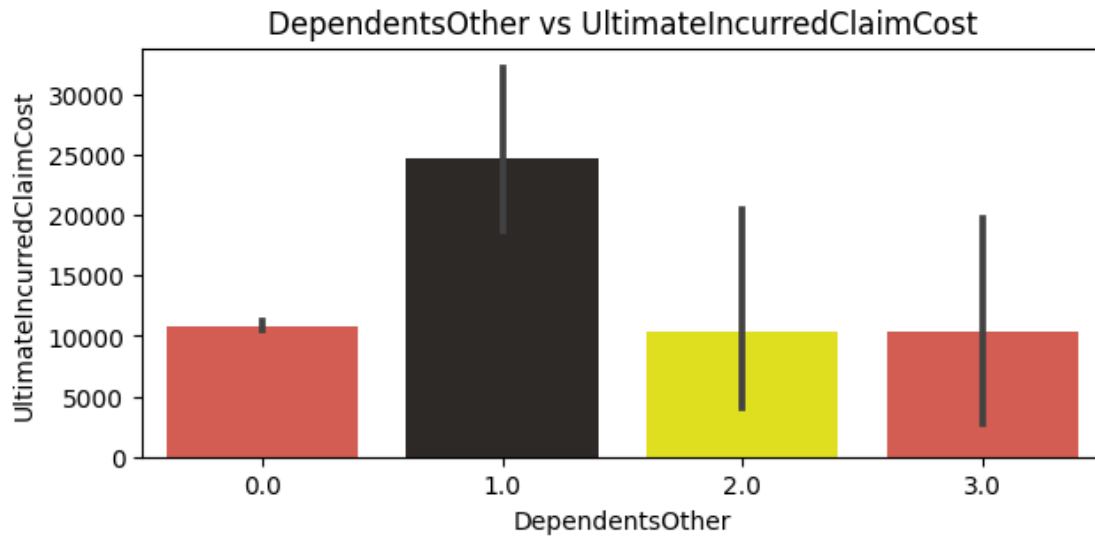
```
[42]: Text(0.5, 1.0, 'DependentChildren vs UltimateIncurredClaimCost')
```



- People who have more children as dependents got more insurance payments from the insurance company.

```
[43]: plt.figure(figsize=(7,3))
sns.barplot(x='DependentsOther',y='UltimateIncurredClaimCost',data=df,
           palette=colors)
plt.title('DependentsOther vs UltimateIncurredClaimCost')
```

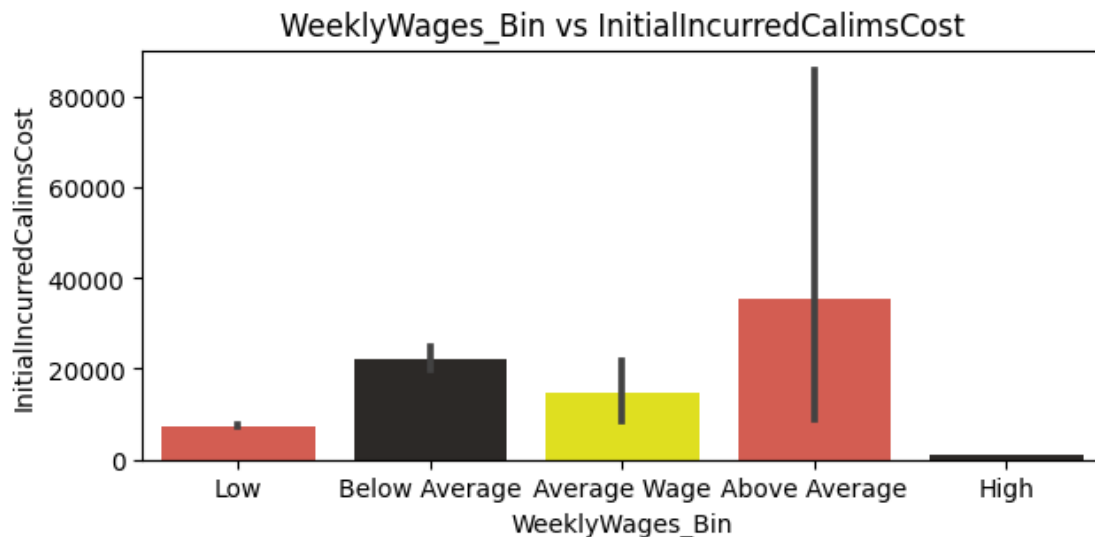
```
[43]: Text(0.5, 1.0, 'DependentsOther vs UltimateIncurredClaimCost')
```



- People having 1 dependent person other than children has highest ultimate claimed cost.

```
[44]: plt.figure(figsize=(7,3))
sns.
    ↳ barplot(x='WeeklyWages_Bin',y='InitialIncurredCalimsCost',data=df,palette=colors)
plt.title('WeeklyWages_Bin vs InitialIncurredCalimsCost')
```

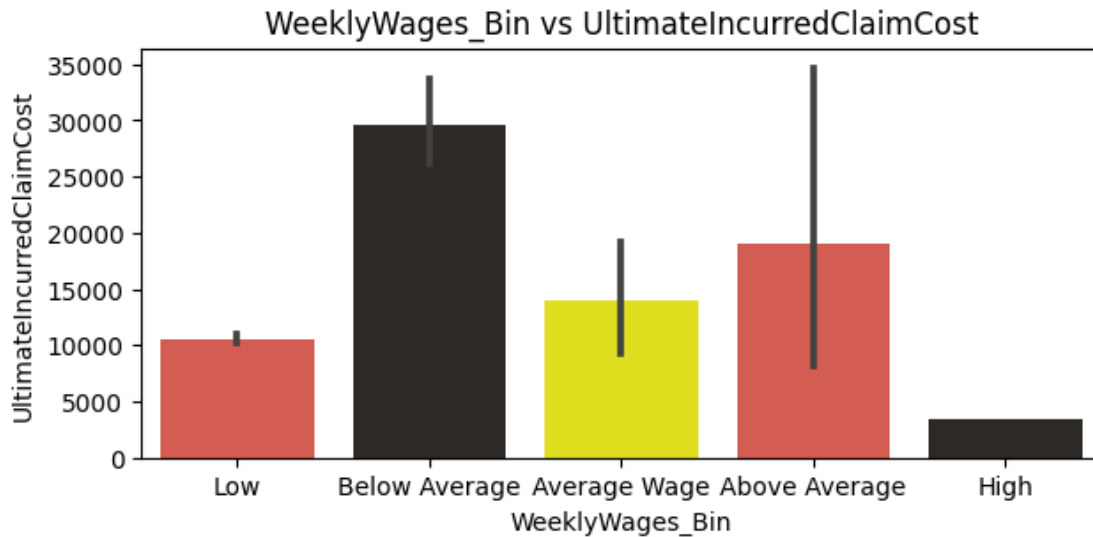
```
[44]: Text(0.5, 1.0, 'WeeklyWages_Bin vs InitialIncurredCalimsCost')
```



- People whose wages are above average claimed for more claim cost.

```
[45]: plt.figure(figsize=(7,3))
sns.
    ↳ barplot(x='WeeklyWages_Bin',y='UltimateIncurredClaimCost',data=df,palette=colors)
plt.title('WeeklyWages_Bin vs UltimateIncurredClaimCost')
```

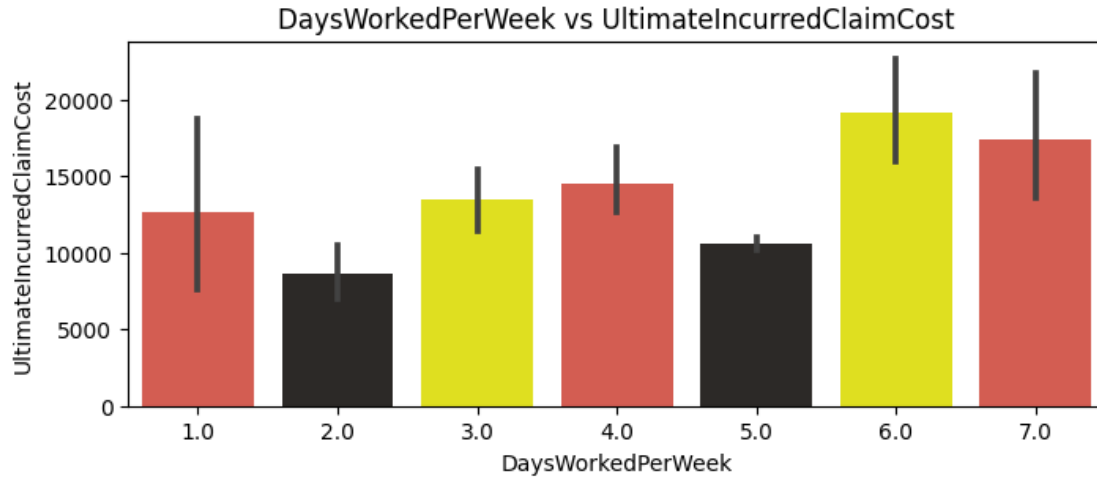
```
[45]: Text(0.5, 1.0, 'WeeklyWages_Bin vs UltimateIncurredClaimCost')
```



- People whose wages are below average and above average got more total claims payments by the insurance company.

```
[46]: plt.figure(figsize=(8,3))
sns.
    ↳ barplot(x='DaysWorkedPerWeek',y='UltimateIncurredClaimCost',data=df,palette=colors)
plt.title('DaysWorkedPerWeek vs UltimateIncurredClaimCost')
```

```
[46]: Text(0.5, 1.0, 'DaysWorkedPerWeek vs UltimateIncurredClaimCost')
```

- People who worked 6-7 days per week got more total claims payments by the insurance company.

5 Feature Engineering

```
[47]: df['YearOfAccident'] = pd.DatetimeIndex(df['DateTimeOfAccident']).year
df['MonthOfAccident'] = pd.DatetimeIndex(df['DateTimeOfAccident']).month
df['DayOfAccident'] = pd.DatetimeIndex(df['DateTimeOfAccident']).day
df['HourOfAccident'] = pd.DatetimeIndex(df['DateTimeOfAccident']).hour
df['YearReported'] = pd.DatetimeIndex(df['DateReported']).year

df['ReportDelayInDays'] = pd.DatetimeIndex(df['DateReported']).date - pd.
    ↳DatetimeIndex(df['DateTimeOfAccident']).date
df['ReportDelayInDays'] = (df['ReportDelayInDays'] / np.timedelta64(1, 'D')).
    ↳astype(int)
df['ReportDelayInWeeks'] = np.floor(df['ReportDelayInDays'] / 7.).astype(int)
df['ReportDelayInWeeks'] = np.clip(df['ReportDelayInWeeks'], a_max=55,
    ↳a_min=None)
```

```
[48]: df.reset_index(drop=True)
```

```
[48]:
```

	ClaimNumber	DateTimeOfAccident	DateReported	Age	\
0	WC8205482	2002-04-09 07:00:00+00:00	2002-07-05 00:00:00+00:00	48	
1	WC6922469	1999-01-07 11:00:00+00:00	1999-01-20 00:00:00+00:00	43	
2	WC5442654	1996-03-25 00:00:00+00:00	1996-04-14 00:00:00+00:00	30	
3	WC9796897	2005-06-22 13:00:00+00:00	2005-07-22 00:00:00+00:00	41	
4	WC2603726	1990-08-29 08:00:00+00:00	1990-09-27 00:00:00+00:00	36	
...	
36171	WC5624756	1996-05-29 09:00:00+00:00	1996-06-27 00:00:00+00:00	20	

36172	WC8516685	2002-10-08	08:00:00+00:00	2003-02-07	00:00:00+00:00	35
36173	WC6891668	1999-09-22	09:00:00+00:00	1999-11-11	00:00:00+00:00	52
36174	WC4287842	1993-02-05	06:00:00+00:00	1993-03-19	00:00:00+00:00	28
36175	WC6368063	1998-03-06	10:00:00+00:00	1998-04-09	00:00:00+00:00	29

	Gender	MaritalStatus	DependentChildren	DependentsOther	WeeklyWages	\
0	M	M	0	0.0	500.00	
1	F	M	0	0.0	509.34	
2	M	U	0	0.0	709.10	
3	M	S	0	0.0	555.46	
4	M	M	0	0.0	377.10	
...	
36171	F	S	0	0.0	344.16	
36172	M	M	0	0.0	1668.83	
36173	F	M	0	0.0	204.87	
36174	M	M	0	0.0	730.87	
36175	M	S	0	0.0	200.00	

	PartTimeFullTime	...	UltimateIncurredClaimCost	Age_Bin	\
0	F	...	4748.203388	Middle-Age	
1	F	...	6326.285819	Middle-Age	
2	F	...	2293.949087	Middle-Age	
3	F	...	17786.487170	Middle-Age	
4	F	...	4014.002925	Middle-Age	
...	
36171	F	...	1343.054886	Young	
36172	F	...	172876.632600	Middle-Age	
36173	P	...	632.281472	Old	
36174	F	...	6714.495760	Middle-Age	
36175	F	...	2588.845117	Middle-Age	

	WeeklyWages_Bin	YearOfAccident	MonthOfAccident	DayOfAccident	\
0	Low	2002	4	9	
1	Low	1999	1	7	
2	Low	1996	3	25	
3	Low	2005	6	22	
4	Low	1990	8	29	
...	
36171	Low	1996	5	29	
36172	Below Average	2002	10	8	
36173	Low	1999	9	22	
36174	Low	1993	2	5	
36175	Low	1998	3	6	

	HourOfAccident	YearReported	ReportDelayInDays	ReportDelayInWeeks
0	7	2002	87	12
1	11	1999	13	1

2	0	1996	20	2
3	13	2005	30	4
4	8	1990	29	4
...
36171	9	1996	29	4
36172	8	2003	122	17
36173	9	1999	50	7
36174	6	1993	42	6
36175	10	1998	34	4

[36176 rows x 24 columns]

```
[49]: numerical_features = [c for c in df.columns if df[c].dtype in ['float64',
    ↪ 'int64', 'int32'] if c not in ['Acc_Day', 'Acc_Month', 'Acc_Year']]
categorical_features = [c for c in df.columns if c not in numerical_features]
```

```
[50]: numerical_features
```

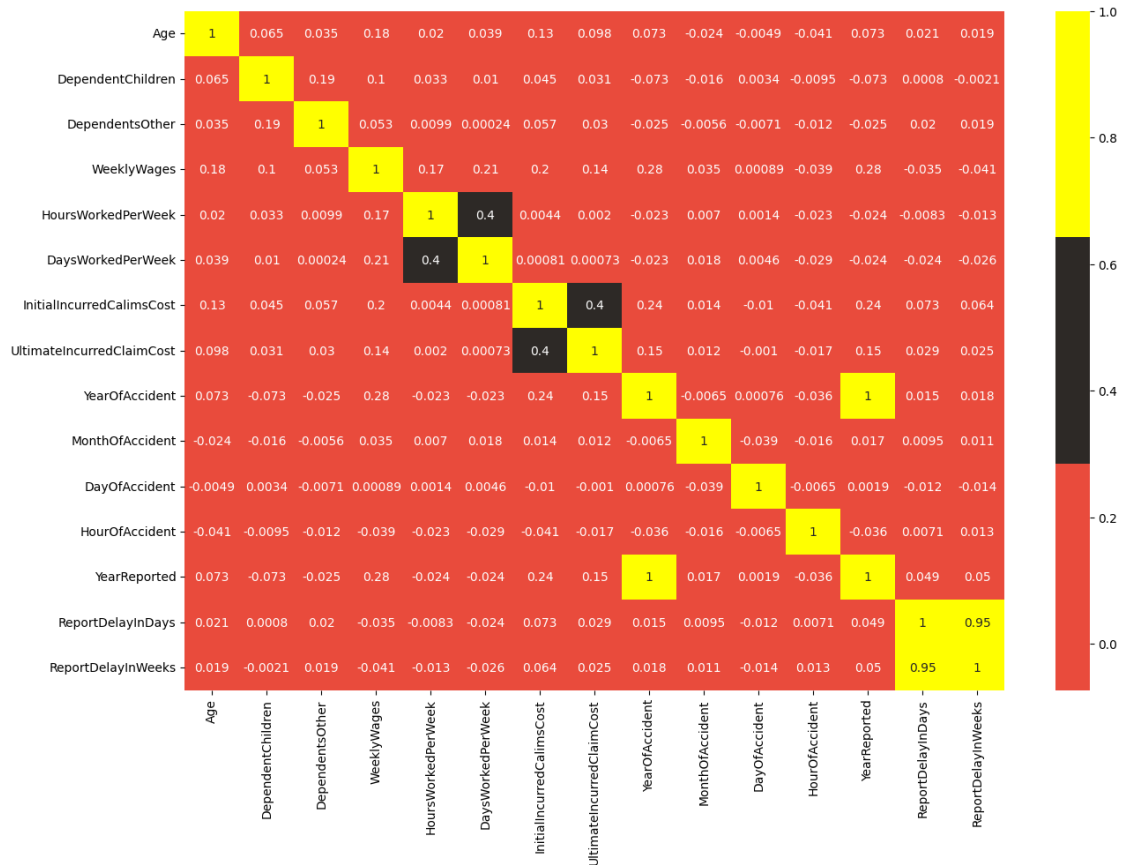
```
[50]: ['Age',
      'DependentChildren',
      'DependentsOther',
      'WeeklyWages',
      'HoursWorkedPerWeek',
      'DaysWorkedPerWeek',
      'InitialIncurredCalimsCost',
      'UltimateIncurredClaimCost',
      'YearOfAccident',
      'MonthOfAccident',
      'DayOfAccident',
      'HourOfAccident',
      'YearReported',
      'ReportDelayInDays',
      'ReportDelayInWeeks']
```

```
[51]: categorical_features
```

```
[51]: ['ClaimNumber',
      'DateTimeOfAccident',
      'DateReported',
      'Gender',
      'MaritalStatus',
      'PartTimeFullTime',
      'ClaimDescription',
      'Age_Bin',
      'WeeklyWages_Bin']
```

5.1 Correlation

```
[52]: plt.figure(figsize=(15,10))
sns.heatmap(df[numerical_features].corr(),annot=True , cmap=colors)
plt.show()
```



- ‘DaysWorkedPerWeek’ and ‘HoursWorkedPerWeek’ have a correlation of 0.4, which makes sense as they both refer to the time worked per week in different metrics. We will use either one of the 2 in our final model as they both provide the same information.
- ‘InitialIncurredCalimsCost’ and ‘UltimateIncurredClaimCost’ have a correlation of 0.4, which is good as we are trying to predict the UltimateIncurredClaimCost and the Ultimate claim cost will be dependent on the Initial claim cost.
- Log transformation is applied on ‘InitialIncurredCalimsCost’ and ‘UltimateIncurredClaimCost’ to transform their skewed distributions to approximately normal. This makes the interpretation much easier

```
[53]: plt.figure(figsize = (10,5))
plt.subplot(1, 2, 1)
sns.distplot(df["UltimateIncurredClaimCost"])
```

```

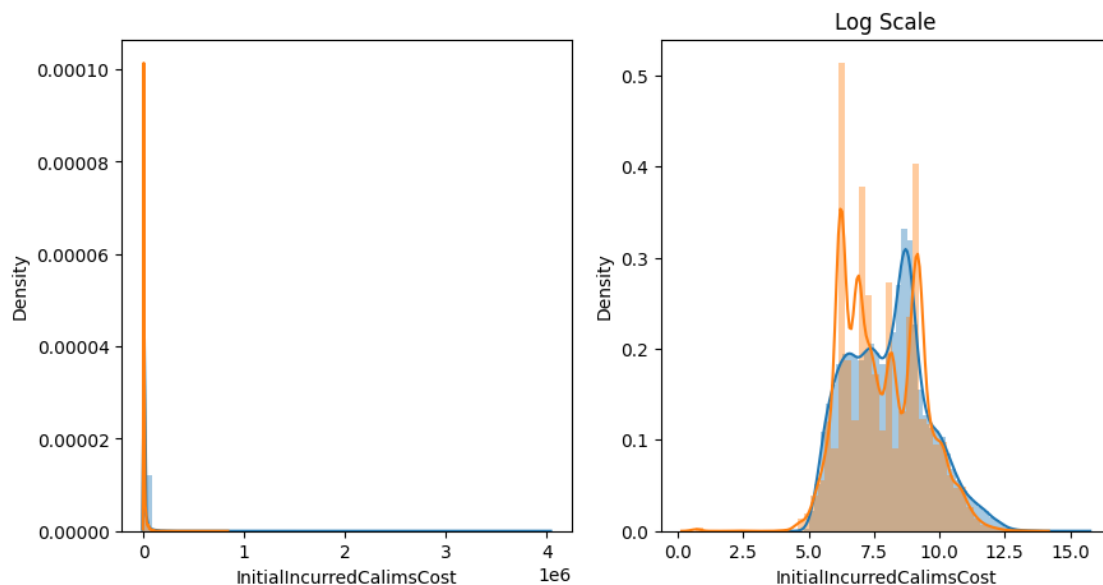
sns.distplot(df['InitialIncurredCalimsCost'])
plt.subplot(1, 2, 2)
plt.title('Log Scale')
sns.distplot(np.log1p(df['UltimateIncurredClaimCost']))
sns.distplot(np.log1p(df['InitialIncurredCalimsCost']))

```

```

[53]: <Axes: title={ 'center': 'Log Scale'}, xlabel='InitialIncurredCalimsCost',
      ylabel='Density'>

```



- The above distribution plot after the log transformation can now be easily understood and interpreted. As the Initial estimate by the insurer of the claim cost increases the total claim payments to be made by the insurance company rises

5.2 Outlier Analysis and Treatment

```

[54]: def count_outliers(df):
      # List of numeric columns
      num_cols = df.select_dtypes(include=['float64', 'int64', 'int32']).columns

      # Calculating quartiles and IQR
      Q1 = df[num_cols].quantile(0.25)
      Q3 = df[num_cols].quantile(0.75)
      IQR = Q3 - Q1

      print('Count of all outliers:\n')
      print((((df[num_cols] < (Q1 - 1.5 * IQR)) | (df[num_cols] > (Q3 + 1.5 * IQR)))
      ↪ .sum()))

```

```
count_outliers(df)
```

Count of all outliers:

Age	15
DependentChildren	2278
DependentsOther	318
WeeklyWages	987
HoursWorkedPerWeek	4929
DaysWorkedPerWeek	3180
InitialIncurredCalimsCost	2889
UltimateIncurredClaimCost	4525
YearOfAccident	0
MonthOfAccident	0
DayOfAccident	0
HourOfAccident	916
YearReported	0
ReportDelayInDays	3336
ReportDelayInWeeks	4273

dtype: int64

Insight:

- 'Age' : has low number of outliers can be ignored
- 'DependentChildren' and 'DependentsOther': the numbers of outliers are more but removing records on the basis of these columns may result in deletion of valuable records.
- 'WeeklyWages' : outlier treatment to be performed.
- 'DaysWorkedPerWeek' and 'HoursWorkedPerWeek': removing records from one of them should handle the other column as well because they are related.
- 'InitialIncurredCalimsCost' : outlier treatment to be performed.
- 'UltimateIncurredClaimCost' : outlier treatment to be performed.

```
[55]: # Function to calculate upper and lower bounds for outliers
def outliers_limits(df, feature):
    Q1 = df[feature].quantile(0.25)
    Q3 = df[feature].quantile(0.75)
    IQR = Q3 - Q1
    upper_limit = Q3 + 1.5 * IQR
    lower_limit = Q1 - 1.5 * IQR
    return upper_limit, lower_limit

# Function to remove outliers based on the calculated limits
def removal(df, feature, upper, lower):
    new_df = df[(df[feature] > lower) & (df[feature] < upper)]
    return new_df
```

```

# Example usage:
feature = [
    'WeeklyWages', 'InitialIncurredCalimsCost', 'UltimateIncurredClaimCost', 'HourOfAccident', 'Re
upper_limit, lower_limit = outliers_limits(df, feature)
cleaned_df = removal(df, feature, upper_limit, lower_limit)

```

```

[56]: def plot_outliers(df, feature, cleaned_df, upper, lower):
    fig, axes = plt.subplots(nrows=1, ncols=len(feature), figsize=(25, 5))

    for i, feat in enumerate(feature):
        # Before removal
        df.boxplot(column=feat, ax=axes[i], color='black')
        axes[i].set_title(f'Before outlier removal - {feat}')
        axes[i].set_ylabel(f'{feat} value')
        axes[i].axhline(y=upper[feat], color='r', linestyle='--', label='Upper_
↳limit')
        axes[i].axhline(y=lower[feat], color='g', linestyle='--', label='Lower_
↳limit')
        axes[i].legend()

    fig.tight_layout()
    plt.show()

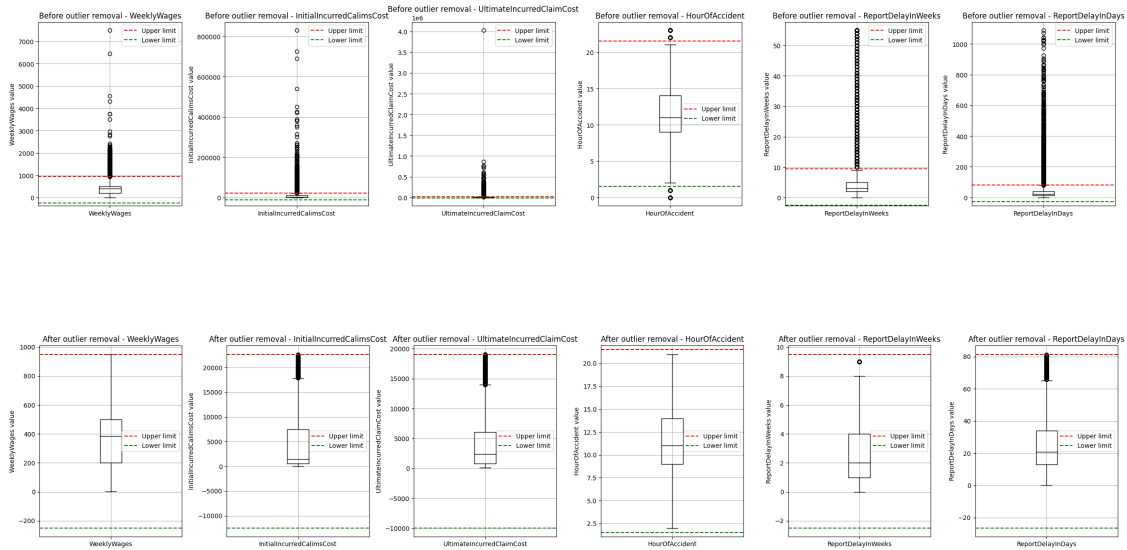
    fig, axes = plt.subplots(nrows=1, ncols=len(feature), figsize=(25, 5))

    for i, feat in enumerate(feature):
        # After removal
        cleaned_df.boxplot(column=feat, ax=axes[i], color='black')
        axes[i].set_title(f'After outlier removal - {feat}')
        axes[i].set_ylabel(f'{feat} value')
        axes[i].axhline(y=upper[feat], color='r', linestyle='--', label='Upper_
↳limit')
        axes[i].axhline(y=lower[feat], color='g', linestyle='--', label='Lower_
↳limit')
        axes[i].legend()

    fig.tight_layout()
    plt.show()

plot_outliers(df, feature, cleaned_df, upper_limit, lower_limit)

```



```
[57]: #dropping unnecessary columns
df=df.
      ↪drop(columns=['WeeklyWages_Bin','Age_Bin','ClaimNumber','DateTimeOfAccident','DateReported'])
df.head()
```

```
[57]:
```

	Age	Gender	MaritalStatus	DependentChildren	DependentsOther	WeeklyWages \
1	48	M	M	0	0.0	500.00
2	43	F	M	0	0.0	509.34
3	30	M	U	0	0.0	709.10
4	41	M	S	0	0.0	555.46
5	36	M	M	0	0.0	377.10

	PartTimeFullTime	HoursWorkedPerWeek	DaysWorkedPerWeek \
1	F	38.0	5.0
2	F	37.5	5.0
3	F	38.0	5.0
4	F	38.0	5.0
5	F	38.0	5.0

	InitialIncurredCalimsCost	UltimateIncurredClaimCost	YearOfAccident \
1	1500	4748.203388	2002
2	5500	6326.285819	1999
3	1700	2293.949087	1996
4	15000	17786.487170	2005
5	2800	4014.002925	1990

	MonthOfAccident	DayOfAccident	HourOfAccident	YearReported \
1	4	9	7	2002

2	1	7	11	1999
3	3	25	0	1996
4	6	22	13	2005
5	8	29	8	1990

	ReportDelayInDays	ReportDelayInWeeks
1	87	12
2	13	1
3	20	2
4	30	4
5	29	4

6 Data Loading for Test Data

```
[58]: df_test = pd.read_csv('Test_SJC.csv')
df_test.head()
```

```
[58]: ClaimNumber    DateTimeOfAccident    DateReported    Age Gender \
0    WC8476284    2002-04-19T16:00:00Z    2002-05-13T00:00:00Z    38    M
1    WC2445024    1989-09-26T08:00:00Z    1989-10-14T00:00:00Z    38    F
2    WC4566945    1994-05-02T13:00:00Z    1994-05-17T00:00:00Z    24    M
3    WC9911299    2005-11-26T06:00:00Z    2006-01-07T00:00:00Z    21    M
4    WC9066190    2003-03-12T13:00:00Z    2003-04-10T00:00:00Z    32    M
```

	MaritalStatus	DependentChildren	DependentsOther	WeeklyWages	\
0	M	0	0	500.00	
1	M	0	0	350.00	
2	S	0	0	487.50	
3	S	0	0	431.62	
4	M	3	0	480.50	

	PartTimeFullTime	HoursWorkedPerWeek	DaysWorkedPerWeek	\
0	F	40.00	5	
1	P	29.75	4	
2	F	38.00	5	
3	F	40.00	5	
4	F	38.00	5	

	ClaimDescription	\
0	STRUCK VALVES ABRASIONS LEFT LEG LACERATED LEF...	
1	LIFTING PATIENT PAIN IN LOWER BACK LEG	
2	LIFTING BOXES LOWER BACK BACK INJURY	
3	STRUCK LADDER BRUISED RIGHT KNEE MUSCLE RIGHT	
4	FELL OFF LADDER FRACTURE RIGHT WRIST	

InitialIncurredCalimsCost

```

0          1000
1          3500
2          7500
3          1000
4         111077

```

```
[59]: df_test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17824 entries, 0 to 17823
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ClaimNumber                          17824 non-null  object
1   DateTimeOfAccident                  17824 non-null  object
2   DateReported                        17824 non-null  object
3   Age                                 17824 non-null  int64
4   Gender                              17824 non-null  object
5   MaritalStatus                       17817 non-null  object
6   DependentChildren                  17824 non-null  int64
7   DependentsOther                    17824 non-null  int64
8   WeeklyWages                        17824 non-null  float64
9   PartTimeFullTime                   17824 non-null  object
10  HoursWorkedPerWeek                 17824 non-null  float64
11  DaysWorkedPerWeek                  17824 non-null  int64
12  ClaimDescription                   17824 non-null  object
13  InitialIncurredCalimsCost          17824 non-null  int64
dtypes: float64(2), int64(5), object(7)
memory usage: 1.9+ MB

```

```
[60]: df_test.isnull().sum()
```

```

[60]: ClaimNumber          0
      DateTimeOfAccident    0
      DateReported         0
      Age                  0
      Gender               0
      MaritalStatus        7
      DependentChildren    0
      DependentsOther      0
      WeeklyWages          0
      PartTimeFullTime     0
      HoursWorkedPerWeek   0
      DaysWorkedPerWeek    0
      ClaimDescription     0
      InitialIncurredCalimsCost  0
      dtype: int64

```

```
[61]: #handling missing values in test data
df_test['MaritalStatus'] = df_test['MaritalStatus'].fillna('U')
```

```
[62]: df_test.isnull().sum()
```

```
[62]: ClaimNumber          0
      DateTimeOfAccident    0
      DateReported          0
      Age                  0
      Gender               0
      MaritalStatus         0
      DependentChildren     0
      DependentsOther       0
      WeeklyWages           0
      PartTimeFullTime      0
      HoursWorkedPerWeek    0
      DaysWorkedPerWeek     0
      ClaimDescription      0
      InitialIncurredCalimsCost  0
      dtype: int64
```

```
[63]: # Type casting : Converting the dtype of 'DateTimeOfAccident' and
      ↪ 'DateReported' from object to datetime64

df_test['DateTimeOfAccident']=pd.to_datetime(df_test['DateTimeOfAccident'])
df_test['DateReported']=pd.to_datetime(df_test['DateReported'])
```

```
[64]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17824 entries, 0 to 17823
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ClaimNumber           17824 non-null  object
1   DateTimeOfAccident    17824 non-null  datetime64[ns, UTC]
2   DateReported          17824 non-null  datetime64[ns, UTC]
3   Age                   17824 non-null  int64
4   Gender                17824 non-null  object
5   MaritalStatus         17824 non-null  object
6   DependentChildren     17824 non-null  int64
7   DependentsOther       17824 non-null  int64
8   WeeklyWages           17824 non-null  float64
9   PartTimeFullTime      17824 non-null  object
10  HoursWorkedPerWeek    17824 non-null  float64
11  DaysWorkedPerWeek     17824 non-null  int64
12  ClaimDescription      17824 non-null  object
13  InitialIncurredCalimsCost 17824 non-null  int64
```

```
dtypes: datetime64[ns, UTC](2), float64(2), int64(5), object(5)
memory usage: 1.9+ MB
```

[65]: *#feature engineering for test data*

```
df_test['YearOfAccident'] = pd.DatetimeIndex(df_test['DateTimeOfAccident']).
    ↳year
df_test['MonthOfAccident'] = pd.DatetimeIndex(df_test['DateTimeOfAccident']).
    ↳month
df_test['DayOfAccident'] = pd.DatetimeIndex(df_test['DateTimeOfAccident']).day
df_test['HourOfAccident'] = pd.DatetimeIndex(df_test['DateTimeOfAccident']).
    ↳hour
df_test['YearReported'] = pd.DatetimeIndex(df_test['DateReported']).year

df_test['ReportDelayInDays'] = pd.DatetimeIndex(df_test['DateReported']).date -
    ↳pd.DatetimeIndex(df_test['DateTimeOfAccident']).date
df_test['ReportDelayInDays'] = (df_test['ReportDelayInDays'] / np.
    ↳timedelta64(1, 'D')).astype(int)
df_test['ReportDelayInWeeks'] = np.floor(df_test['ReportDelayInDays'] / 7.).
    ↳astype(int)
df_test['ReportDelayInWeeks'] = np.clip(df_test['ReportDelayInWeeks'],
    ↳a_max=55, a_min=None)
```

[66]: df_test.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17824 entries, 0 to 17823
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ClaimNumber                          17824 non-null  object
1   DateTimeOfAccident                   17824 non-null  datetime64[ns, UTC]
2   DateReported                        17824 non-null  datetime64[ns, UTC]
3   Age                                  17824 non-null  int64
4   Gender                               17824 non-null  object
5   MaritalStatus                       17824 non-null  object
6   DependentChildren                   17824 non-null  int64
7   DependentsOther                     17824 non-null  int64
8   WeeklyWages                         17824 non-null  float64
9   PartTimeFullTime                    17824 non-null  object
10  HoursWorkedPerWeek                  17824 non-null  float64
11  DaysWorkedPerWeek                   17824 non-null  int64
12  ClaimDescription                     17824 non-null  object
13  InitialIncurredCalimsCost            17824 non-null  int64
14  YearOfAccident                       17824 non-null  int32
15  MonthOfAccident                     17824 non-null  int32
16  DayOfAccident                       17824 non-null  int32
17  HourOfAccident                       17824 non-null  int32
```

```

18 YearReported          17824 non-null  int32
19 ReportDelayInDays     17824 non-null  int32
20 ReportDelayInWeeks    17824 non-null  int32
dtypes: datetime64[ns, UTC](2), float64(2), int32(7), int64(5), object(5)
memory usage: 2.4+ MB

```

```

[67]: #dropping unimportant features
df_test = df_test.drop(['ClaimNumber', 'DateTimeOfAccident', 'DateReported',
↳ 'ClaimDescription'], axis=1)

```

7 Label Encoding

- Converting text in to numerical values

```

[68]: gender_label = {'M':1, 'F':2, 'U': 3}
marital_label = {'M':1, 'S':2, 'U':3}
partTime_label = {'F':1, 'P':2}

df['Gender'] = df['Gender'].map(gender_label)
df['MaritalStatus'] = df['MaritalStatus'].map(marital_label)
df['PartTimeFullTime'] = df['PartTimeFullTime'].map(partTime_label)

```

```

[69]: #for test data

df_test['Gender'] = df_test['Gender'].map(gender_label)
df_test['MaritalStatus'] = df_test['MaritalStatus'].map(marital_label)
df_test['PartTimeFullTime'] = df_test['PartTimeFullTime'].map(partTime_label)

```

```

[70]: df.head()

```

```

[70]:   Age  Gender  MaritalStatus  DependentChildren  DependentsOther  \
1   48      1              1              0          0.0
2   43      2              1              0          0.0
3   30      1              3              0          0.0
4   41      1              2              0          0.0
5   36      1              1              0          0.0

   WeeklyWages  PartTimeFullTime  HoursWorkedPerWeek  DaysWorkedPerWeek  \
1       500.00              1          38.0          5.0
2       509.34              1          37.5          5.0
3       709.10              1          38.0          5.0
4       555.46              1          38.0          5.0
5       377.10              1          38.0          5.0

   InitialIncurredCalimsCost  UltimateIncurredClaimCost  YearOfAccident  \
1              1500          4748.203388          2002
2              5500          6326.285819          1999

```

3	1700	2293.949087	1996
4	15000	17786.487170	2005
5	2800	4014.002925	1990

	MonthOfAccident	DayOfAccident	HourOfAccident	YearReported	\
1	4	9	7	2002	
2	1	7	11	1999	
3	3	25	0	1996	
4	6	22	13	2005	
5	8	29	8	1990	

	ReportDelayInDays	ReportDelayInWeeks
1	87	12
2	13	1
3	20	2
4	30	4
5	29	4

```
[71]: df_test.head()
```

```
[71]:
```

	Age	Gender	MaritalStatus	DependentChildren	DependentsOther	\
0	38	1	1	0	0	
1	38	2	1	0	0	
2	24	1	2	0	0	
3	21	1	2	0	0	
4	32	1	1	3	0	

	WeeklyWages	PartTimeFullTime	HoursWorkedPerWeek	DaysWorkedPerWeek	\
0	500.00	1	40.00	5	
1	350.00	2	29.75	4	
2	487.50	1	38.00	5	
3	431.62	1	40.00	5	
4	480.50	1	38.00	5	

	InitialIncurredCalimsCost	YearOfAccident	MonthOfAccident	DayOfAccident	\
0	1000	2002	4	19	
1	3500	1989	9	26	
2	7500	1994	5	2	
3	1000	2005	11	26	
4	111077	2003	3	12	

	HourOfAccident	YearReported	ReportDelayInDays	ReportDelayInWeeks
0	16	2002	24	3
1	8	1989	18	2
2	13	1994	15	2
3	6	2006	42	6
4	13	2003	29	4

8 Data Normalization

- Normalizing features ensures that all features have a similar scale
- Before normalization we should do train test split so that we can avoid data leakage issues.
- In order to avoid Data Leakage, it is advised to use train-test-split before any transformations. Execute the transformations according to the training data for the training as well as test data.

```
[72]: from sklearn.model_selection import train_test_split

train , test = train_test_split(df, test_size = 0.3)

x_train = train.drop('UltimateIncurredClaimCost', axis=1)
y_train = train['UltimateIncurredClaimCost']

x_test = test.drop('UltimateIncurredClaimCost', axis = 1)
y_test = test['UltimateIncurredClaimCost']
```

```
[73]: #for train data
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler(feature_range=(0, 1))

x_train_scaled = scaler.fit_transform(x_train)
x_train = pd.DataFrame(x_train_scaled)

x_test_scaled = scaler.transform(x_test)
x_test = pd.DataFrame(x_test_scaled)
```

```
[74]: #for test data

df_test_scaled=scaler.fit_transform(df_test)
(df_test-df_test.min())/(df_test.max()-df_test)
scaler.fit(df_test)
scaler.transform(df_test)
```

```
[74]: array([[0.36764706, 0.          , 0.          , ..., 0.77777778, 0.02191781,
              0.05454545],
              [0.36764706, 0.5        , 0.          , ..., 0.05555556, 0.01643836,
              0.03636364],
              [0.16176471, 0.          , 0.5        , ..., 0.33333333, 0.01369863,
              0.03636364],
              ...,
              [0.08823529, 0.          , 0.5        , ..., 0.22222222, 0.0173516 ,
              0.03636364],
              [0.16176471, 0.          , 0.5        , ..., 0.38888889, 0.01917808,
              0.05454545],
              [0.13235294, 0.          , 0.5        , ..., 0.11111111, 0.00913242,
              0.01818182]])
```

9 Modeling

9.1 Linear regression

It is a machine learning algorithm based on supervised learning. It performs a regression task. Through linear regression we can find out the linear relationship between the target and the explanatory variables.

```
[75]: from sklearn.metrics import mean_squared_error
      from sklearn.linear_model import LinearRegression
      from sklearn.metrics import r2_score
```

```
[76]: reg = LinearRegression()
      reg.fit(x_train, y_train)
```

```
[76]: LinearRegression()
```

```
[77]: # predicting the test set results
      y_pred = reg.predict(x_test)

      # Calculating the r2 score
      r2 = r2_score(y_test, y_pred)
      print("r2 score :", r2)
```

```
r2 score : 0.10562321552454579
```

- Root Mean Squared Error (RMSE) is a metric used to evaluate a Regression Model. It tells us how accurate our predictions are and what is the amount of deviation from the actual values.

```
[78]: print((f"Regression RMSE: {np.sqrt(mean_squared_error(y_test, y_pred))}"))
```

```
Regression RMSE: 45410.651266437555
```

```
[79]: reg.score(x_train,y_train)
```

```
[79]: 0.24823914659803303
```

9.2 Random Forest regression

Every decision tree has high variance, but when we combine all of them together in parallel then the resultant variance is low as each decision tree gets perfectly trained on that particular sample data and hence the output doesn't depend on one decision tree but multiple decision trees.

Random Forest is an ensemble technique capable of performing both regression and classification tasks with the use of multiple decision trees and a technique called Bootstrap and Aggregation, commonly known as bagging. The basic idea behind this is to combine multiple decision trees in determining the final output rather than relying on individual decision trees.

```
[80]: from sklearn.ensemble import RandomForestRegressor
```



```
[81]: rf = RandomForestRegressor()  
      rf.fit(x_train, y_train)
```

```
[81]: RandomForestRegressor()
```

```
[82]: y_pred1 = rf.predict(x_test)  
  
      # Calculating the r2 score  
      r2 = r2_score(y_test, y_pred1)  
      print("r2 score :", r2)
```

```
r2 score : 0.11344291638567794
```

```
[83]: print((f"Regression RMSE: {np.sqrt(mean_squared_error(y_test, y_pred1))}"))
```

```
Regression RMSE: 45211.69860023026
```

```
[84]: rf.score(x_train,y_train)
```

```
[84]: 0.8967421622463381
```

9.3 LGBM Regressor

LightGBM (Light Gradient Boosting Machine) Regressor is a type of gradient boosting model specifically designed for efficiency, speed, and high performance on large datasets. It is a tree-based ensemble model that builds multiple decision trees sequentially, where each tree corrects the errors of the previous one.

```
[85]: from lightgbm import LGBMRegressor  
      lgbm =LGBMRegressor()  
  
      lgbm_model = lgbm.fit(x_train, y_train)  
      lg_vpreds = lgbm_model.predict(x_test)  
  
      r2 = r2_score(y_test, lg_vpreds)  
      print("r2 score :", r2)  
      print((f"LGBM RMSE: {np.sqrt(mean_squared_error(y_test, lg_vpreds))}"))
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of  
testing was 0.001179 seconds.
```

```
You can set `force_row_wise=true` to remove the overhead.
```

```
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 1154
```

```
[LightGBM] [Info] Number of data points in the train set: 25323, number of used  
features: 17
```

```
[LightGBM] [Info] Start training from score 10718.897305
```

```
r2 score : 0.14521843950216928
```

```
LGBM RMSE: 44394.07834376454
```

9.4 XG Boost Regression

It is highly optimized and provides fast and accurate predictions for regression tasks. XGBoost uses a more traditional depth-wise tree growth strategy compared to LightGBM's leaf-wise strategy. It supports parallel and distributed computing, making it scalable to large datasets. XGBoost is known for its flexibility and robustness, as it provides a wide range of hyperparameters for fine-tuning and regularization.

```
[86]: from xgboost import XGBRegressor

xgb = XGBRegressor()
xgb_model = xgb.fit(x_train, y_train)
xg_vpreds = xgb_model.predict(x_test)

r2 = r2_score(y_test, lg_vpreds)
print("r2 score :", r2)
print((f"XGBRegressor RMSE: {np.sqrt(mean_squared_error(y_test, xg_vpreds))}"))
```

```
r2 score : 0.14521843950216928
XGBRegressor RMSE: 45382.90029449528
```

- Lower RMSE values indicate a better fit. The RMSE value for LightGBM Regression model is lower as compared to other Regression models. Hence, we go with LightGBM Regression model to predict our Target Variable and further we will do hyperparameter tuning.

9.5 Hyperparameter Tuning for LightGBM regressor

```
[87]: from lightgbm import LGBMRegressor

lgbm =LGBMRegressor(objective = 'regression',
                    num_leaves = 4,
                    learning_rate = 0.01,
                    n_estimators = 10000,
                    max_bin = 200,
                    bagging_fraction = 0.75,
                    bagging_freq = 5,
                    bagging_seed = 7,
                    feature_fraction = 0.2,
                    feature_fraction_seed = 7,
                    verbose = 1,)

lgbm_model = lgbm.fit(x_train, y_train)
lg_vpreds = lgbm_model.predict(x_test)

r2 = r2_score(y_test, lg_vpreds)
```

```
[LightGBM] [Warning] feature_fraction is set=0.2, colsample_bytree=1.0 will be
ignored. Current value: feature_fraction=0.2
[LightGBM] [Warning] bagging_fraction is set=0.75, subsample=1.0 will be
```

ignored. Current value: bagging_fraction=0.75
 [LightGBM] [Warning] bagging_freq is set=5, subsample_freq=0 will be ignored.
 Current value: bagging_freq=5
 [LightGBM] [Warning] feature_fraction is set=0.2, colsample_bytree=1.0 will be ignored. Current value: feature_fraction=0.2
 [LightGBM] [Warning] bagging_fraction is set=0.75, subsample=1.0 will be ignored. Current value: bagging_fraction=0.75
 [LightGBM] [Warning] bagging_freq is set=5, subsample_freq=0 will be ignored. Current value: bagging_freq=5
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000704 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.
 [LightGBM] [Info] Total Bins 975
 [LightGBM] [Info] Number of data points in the train set: 25323, number of used features: 17
 [LightGBM] [Info] Start training from score 10718.897305
 [LightGBM] [Warning] feature_fraction is set=0.2, colsample_bytree=1.0 will be ignored. Current value: feature_fraction=0.2
 [LightGBM] [Warning] bagging_fraction is set=0.75, subsample=1.0 will be ignored. Current value: bagging_fraction=0.75
 [LightGBM] [Warning] bagging_freq is set=5, subsample_freq=0 will be ignored. Current value: bagging_freq=5

```
[88]: print("r2 score :", r2)
      print((f"LGBM RMSE: {np.sqrt(mean_squared_error(y_test, lg_vpreds))}"))
```

r2 score : 0.13021427499593796
 LGBM RMSE: 44782.01277618353

- Before hyperparameter -> 44394.07834376454
- After hyperparameter -> 44782.01277618353

10 Result

- Let's predict the target variable

```
[89]: lg_vpreds = lgbm_model.predict(df_test_scaled)
```

[LightGBM] [Warning] feature_fraction is set=0.2, colsample_bytree=1.0 will be ignored. Current value: feature_fraction=0.2
 [LightGBM] [Warning] bagging_fraction is set=0.75, subsample=1.0 will be ignored. Current value: bagging_fraction=0.75
 [LightGBM] [Warning] bagging_freq is set=5, subsample_freq=0 will be ignored. Current value: bagging_freq=5

```
[90]: df_test = df_test
      df_test['UltimateIncurredClaimCost'] = lg_vpreds
```

```
[91]: df_test.head()
```

```
[91]:
```

	Age	Gender	MaritalStatus	DependentChildren	DependentsOther	\
0	38	1	1	0	0	
1	38	2	1	0	0	
2	24	1	2	0	0	
3	21	1	2	0	0	
4	32	1	1	3	0	

	WeeklyWages	PartTimeFullTime	HoursWorkedPerWeek	DaysWorkedPerWeek	\
0	500.00	1	40.00	5	
1	350.00	2	29.75	4	
2	487.50	1	38.00	5	
3	431.62	1	40.00	5	
4	480.50	1	38.00	5	

	InitialIncurredCalimsCost	YearOfAccident	MonthOfAccident	DayOfAccident	\
0	1000	2002	4	19	
1	3500	1989	9	26	
2	7500	1994	5	2	
3	1000	2005	11	26	
4	111077	2003	3	12	

	HourOfAccident	YearReported	ReportDelayInDays	ReportDelayInWeeks	\
0	16	2002	24	3	
1	8	1989	18	2	
2	13	1994	15	2	
3	6	2006	42	6	
4	13	2003	29	4	

	UltimateIncurredClaimCost
0	5240.234310
1	8385.208148
2	20491.982849
3	-2365.767570
4	84406.131682

```
[ ]:
```